

Implementation of a Benchmark Suite for Strymon

Nicolas Hafner

ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

inf |

Informatik
Computer Science



Motivation

- Performance is important for systems
- Measuring performance can be tough
- Would like to compare to competitors
- Industry benchmarks need to be evaluated

Other Works

- Investigated 9 papers for other systems
- Often very simple benchmarks like Word Count
- Code and data often not published
- No paper used a standardised benchmark

Timely

- Streaming processing based on data parallelism
- Computations expressed via data flow graphs
- Each physical worker runs the entire data flow
- Physical exchange of data handled by Timely
- Data is separated into chunks called epochs
- Computation progress tracked by epoch frontiers

Benchmarks

- We implemented three benchmarks:
 1. Intel's HiBench
 2. Yahoo's Streaming Benchmark
 3. Apache Beam's NEXMark

HiBench

- General Big Data benchmark
- Includes a section for streaming systems
- Only four latency tests:

1. Identity



2. Repartition



3. Word Count



4. Window Reduce





YSB

- Initially developed internally at Yahoo:
- Count ad views for ad campaigns
- Original setup requires Redis table lookup
- Only one, relatively simple data flow:



NEXMark

- Based on a paper by Tucker et al.
- An adaptation of XMark for streaming
- Implements an “auctioning system”
- 13 data flows in total
- Uses filter, map, reduce, join, window, session, partition
- Dataflows for Query 5 and 8:

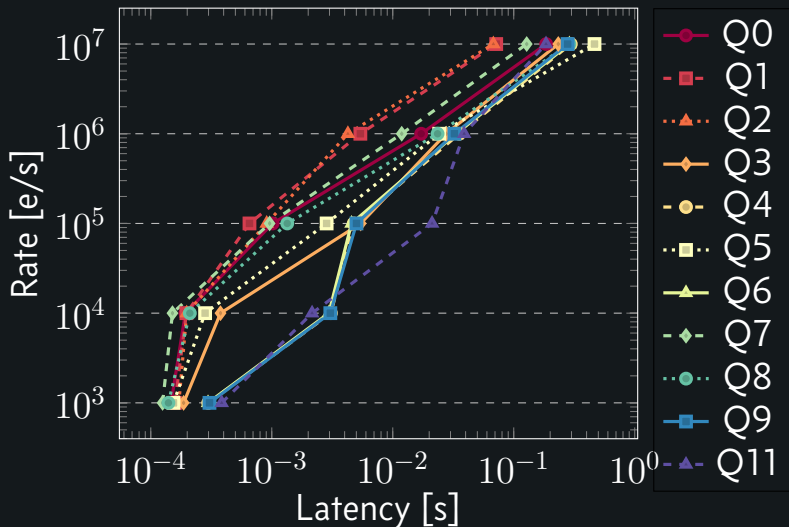


Evaluation

- Run on sgs-r815-03 (32 cores, 2.4GHz)
- Measured closed-loop per-epoch latency
- Data generated directly in memory
- Timely performs very well!


NEXMark Latency Scaling

Median Latency (32 workers)




NEXMark Worker Scaling

Scaling (10'000'000 e/s)



NEXMark CDF

CDF (32 workers, 10'000'000 e/s)



Conclusion

- Timely competes very well with other systems
- Existing benchmarks leave things to be desired
- New benchmarks need:
 - Abstract model definitions for data flows
 - Various short and long data flows
 - Deterministically generated workloads
 - Correctness verification tools

Thanks

Thanks to John Liagouris, Frank McSherry, and the whole DCMModel team!



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich



Informatik
Computer Science

